

# Assignment - AI for business

Utkarsh Thusoo

M20AIE318

## Importing libraries and Loading Data

```
In [1]: import pandas as pd
Data=pd.read_csv("Data_Consumer_ChurnPrediction.csv")
Data
```

```
Out[1]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service
1	5575-GNVDE	Male	0	No	No	34	Yes	No
2	3668-QPYBK	Male	0	No	No	2	Yes	No
3	7795-CFOCW	Male	0	No	No	45	No	No phone service
4	9237-HQITU	Female	0	No	No	2	Yes	No
...	...	...	...	...	...	...	...	...
7038	6840-RESVB	Male	0	Yes	Yes	24	Yes	Yes
7039	2234-XADUH	Female	0	Yes	Yes	72	Yes	Yes
7040	4801-JJAZL	Female	0	Yes	Yes	11	No	No phone service
7041	8361-LTMKD	Male	1	Yes	No	4	Yes	Yes
7042	3186-AJIEK	Male	0	No	No	66	Yes	No

7043 rows × 21 columns

# Preprocessing - Encoding variables

In [2]: `Data.drop('customerID',axis=1,inplace = True) #customerID is not binary and categorical`

In [3]: `Data['TotalCharges'] = pd.to_numeric(Data['TotalCharges'],errors='coerce') #Convert the  
Data['TotalCharges'] = Data['TotalCharges'].fillna(Data['TotalCharges'].median()) #Replace  
  
Data['Churn'] = Data['Churn'].map({'Yes':1, 'No':0})  
Data['gender'] = Data['gender'].map({'Male':1, 'Female':0})  
Data['Partner'] = Data['Partner'].map({'Yes':1, 'No':0}) #  
Data['Dependents'] = Data['Dependents'].map({'Yes':1, 'No':0}) #  
Data['PhoneService'] = Data['PhoneService'].map({'Yes':1, 'No':0}) #  
Data['PaperlessBilling'] = Data['PaperlessBilling'].map({'Yes':1, 'No':0}) #  
  
CatVar = ['MultipleLines','InternetService','OnlineSecurity','OnlineBackup','DevicePr  
Data = pd.get_dummies(Data, columns = CatVar, drop_first=False)  
Data`

Out[3]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	PaperlessBilling	MonthlyCh
--	--------	---------------	---------	------------	--------	--------------	------------------	-----------

0	0	0	1	0	1	0	1	
1	1	0	0	0	34	1	0	
2	1	0	0	0	2	1	1	
3	1	0	0	0	45	0	0	
4	0	0	0	0	2	1	1	
...	...	...	...	...	...	...	...	
7038	1	0	1	1	24	1	1	
7039	0	0	1	1	72	1	1	1
7040	0	0	1	1	11	0	1	
7041	1	1	1	0	4	1	1	
7042	1	0	0	0	66	1	1	1

7043 rows × 41 columns

## Preprocessing - scaling

In [4]: `from sklearn.preprocessing import MinMaxScaler  
sc = MinMaxScaler()  
Data['tenure'] = sc.fit_transform(Data[['tenure']])  
Data['MonthlyCharges'] = sc.fit_transform(Data[['MonthlyCharges']])  
Data['TotalCharges'] = sc.fit_transform(Data[['TotalCharges']])`

# Splitting the data

```
In [5]: from sklearn.model_selection import train_test_split
X = Data.drop('Churn', axis=1)
y = Data['Churn']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=
```

## LogisticRegression

```
In [10]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("accuracy: ", accuracy_score(y_test, y_pred))
print("precision: ", precision_score(y_test, y_pred))
print("recall: ", recall_score(y_test, y_pred))
print("f1_score: ", f1_score(y_test, y_pred, average='weighted'))

accuracy: 0.7988641741599621
precision: 0.6294117647058823
recall: 0.5763016157989228
f1_score: 0.7959305469162359
```

## SVC Classifier

```
In [11]: from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

model = SVC()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("accuracy: ", accuracy_score(y_test, y_pred))
print("precision: ", precision_score(y_test, y_pred))
print("recall: ", recall_score(y_test, y_pred))
print("f1_score: ", f1_score(y_test, y_pred, average='weighted'))

accuracy: 0.7969711310932324
precision: 0.6397379912663755
recall: 0.526032315978456
f1_score: 0.7901995671779648
```

## RandomForestClassifier

```
In [12]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

model = RandomForestClassifier()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

```
print("accuracy: ",accuracy_score(y_test, y_pred))
print("precision: ",precision_score(y_test, y_pred))
print("recall: ",recall_score(y_test, y_pred))
print("f1_score: ",f1_score(y_test, y_pred, average='weighted'))
```

```
accuracy: 0.7823000473260767
precision: 0.6061269146608315
recall: 0.49730700179533216
f1_score: 0.7749541892154461
```

## DecisionTreeClassifier

```
In [13]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

model = DecisionTreeClassifier()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("accuracy: ",accuracy_score(y_test, y_pred))
print("precision: ",precision_score(y_test, y_pred))
print("recall: ",recall_score(y_test, y_pred))
print("f1_score: ",f1_score(y_test, y_pred, average='weighted'))

accuracy: 0.7250354945575012
precision: 0.4807073954983923
recall: 0.5368043087971275
f1_score: 0.7296822584806398
```

## GaussianNB

```
In [14]: from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

model = GaussianNB()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("accuracy: ",accuracy_score(y_test, y_pred))
print("precision: ",precision_score(y_test, y_pred))
print("recall: ",recall_score(y_test, y_pred))
print("f1_score: ",f1_score(y_test, y_pred, average='weighted'))

accuracy: 0.6800757217226692
precision: 0.44495837187789083
recall: 0.8635547576301615
f1_score: 0.6988599745961804
```

## MultinomialNB

```
In [15]: from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

model = GaussianNB()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

```
print("accuracy: ",accuracy_score(y_test, y_pred))
print("precision: ",precision_score(y_test, y_pred))
print("recall: ",recall_score(y_test, y_pred))
print("f1_score: ",f1_score(y_test, y_pred, average='weighted'))
```

```
accuracy: 0.6800757217226692
precision: 0.44495837187789083
recall: 0.8635547576301615
f1_score: 0.6988599745961804
```

## BernoulliNB

```
In [16]: from sklearn.naive_bayes import BernoulliNB
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

model = BernoulliNB()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("accuracy: ",accuracy_score(y_test, y_pred))
print("precision: ",precision_score(y_test, y_pred))
print("recall: ",recall_score(y_test, y_pred))
print("f1_score: ",f1_score(y_test, y_pred, average='weighted'))

accuracy: 0.7094178892569806
precision: 0.47141424272818455
recall: 0.8438061041292639
f1_score: 0.7266304660238987
```

## KNeighborsClassifier

```
In [17]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

model = KNeighborsClassifier()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("accuracy: ",accuracy_score(y_test, y_pred))
print("precision: ",precision_score(y_test, y_pred))
print("recall: ",recall_score(y_test, y_pred))
print("f1_score: ",f1_score(y_test, y_pred, average='weighted'))

accuracy: 0.7657359204921912
precision: 0.5563636363636364
recall: 0.5493716337522442
f1_score: 0.7652581286536987
```

## XGBClassifier

SG Boost

```
In [19]: y_train=y_train.replace("Yes", 1).replace("No", 0)
y_test=y_test.replace("Yes", 1).replace("No", 0)
```

```
In [20]: ##pip install xgboost
from xgboost import XGBClassifier

# declare parameters
params = {
    'objective': 'binary:logistic',
    'max_depth': 4,
    'alpha': 10,
    'learning_rate': 1.0,
    'n_estimators': 100
}
# instantiate the classifier
xgb_clf = XGBClassifier(**params)
# fit the classifier to the training data
xgb_clf.fit(X_train, y_train)
y_pred = xgb_clf.predict(X_test)
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.88	0.86	1556
1	0.61	0.54	0.57	557
accuracy			0.79	2113
macro avg	0.73	0.71	0.72	2113
weighted avg	0.78	0.79	0.78	2113