

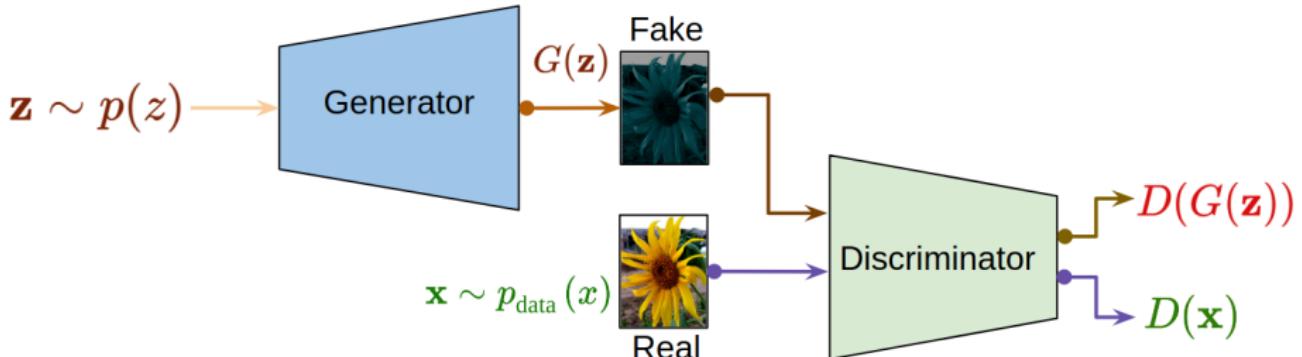
Machine Learning II: Fractal 3

Rajendra Nagar

Assistant Professor
Department of Electrical Engineering
Indian Institute of Technology Jodhpur
<http://home.iitj.ac.in/~rn/>

November 7, 2021

Generative Adversarial Networks: A Minimax Game



- Discriminator wants $D(\mathbf{x}) = 1$ for real samples.
- Discriminator wants $D(G(\mathbf{z})) = 0$ for fake samples.
- Generator wants $D(G(\mathbf{z})) = 1$ for fake samples.

$$\min_G \max_D \left(\underbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})]}_{\text{D wants } D(\mathbf{x}) = 1 \text{ for real point}} + \underbrace{\mathbb{E}_{\mathbf{z} \sim p(z)} [\log(1 - D(G(\mathbf{z})))]}_{\substack{\text{G wants } D(G(\mathbf{z})) = 1 \text{ for fake point} \\ \text{D wants } D(G(\mathbf{z})) = 0 \text{ for fake point}}} \right)$$

Generative Adversarial Networks: Optimality Analysis

Jointly train generator G and discriminator D with a minimax game.

$$\min_G \max_D (\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(z)} [\log(1 - D(G(\mathbf{z})))])$$

We want to verify that global minimum of this game occurs at $p_G = p_{\text{data}}$.

$$\begin{aligned} & \min_G \max_D (\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(z)} [\log(1 - D(G(\mathbf{z})))]) \\ = & \min_G \max_D (\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{x} \sim p_G} [\log(1 - D(\mathbf{x}))]) \\ = & \min_G \max_D \int_{\mathbf{x}} (p_{\text{data}}(\mathbf{x}) [\log D(\mathbf{x})] + p_G(\mathbf{x}) [\log(1 - D(\mathbf{x}))]) d\mathbf{x} \\ = & \min_G \int_{\mathbf{x}} \max_D (p_{\text{data}}(\mathbf{x}) [\log D(\mathbf{x})] + p_G(\mathbf{x}) [\log(1 - D(\mathbf{x}))]) d\mathbf{x} \end{aligned}$$

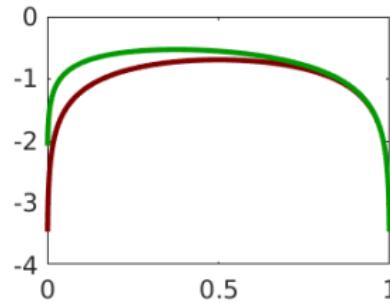
Generative Adversarial Networks: Optimality Analysis

Consider the problem of finding the point of maximum of the function $f(t) = a \log(t) + b \log(1 - t)$. Here $t, a, b \in [0, 1]$.

$$f(t) = a \log(t) + b \log(1 - t)$$

$$\frac{df}{dt} = \frac{a}{t} - \frac{b}{1-t} = 0$$

$$t = \frac{a}{a+b}.$$



$$a = \frac{1}{2}, b = \frac{1}{2}, a = \frac{3}{10}, b = \frac{1}{2}$$

Now, consider the main problem :

$$\max_D (p_{\text{data}}(\mathbf{x})[\log D(\mathbf{x})] + p_G(\mathbf{x})[\log(1 - D(\mathbf{x}))])$$

The optimal discriminator $D^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$

Generative Adversarial Networks

Let us now try to find the optimal generator given the optimal discriminator.

$$\begin{aligned} & \min_G \int_{\mathbf{x}} \max_D (p_{\text{data}}(\mathbf{x}) \log D(\mathbf{x}) + p_G(\mathbf{x}) \log(1 - D(\mathbf{x}))) d\mathbf{x} \\ = & \min_G \int_{\mathbf{x}} (p_{\text{data}}(\mathbf{x}) \log D^*(\mathbf{x}) + p_G(\mathbf{x}) \log(1 - D^*(\mathbf{x}))) d\mathbf{x} \\ = & \min_G \int_{\mathbf{x}} \left(p_{\text{data}} \log \left(\frac{p_{\text{data}}}{p_{\text{data}} + p_G} \right) + p_G \log \left(1 - \frac{p_{\text{data}}}{p_{\text{data}} + p_G} \right) \right) d\mathbf{x} \\ = & \min_G \int_{\mathbf{x}} \left(p_{\text{data}} \log \left(\frac{p_{\text{data}}}{p_{\text{data}} + p_G} \right) + p_G \log \left(\frac{p_G}{p_{\text{data}} + p_G} \right) \right) d\mathbf{x} \\ = & \min_G \int_{\mathbf{x}} \left(p_{\text{data}} \log \left(\frac{2}{2} \frac{p_{\text{data}}}{p_{\text{data}} + p_G} \right) + p_G \log \left(\frac{2}{2} \frac{p_G}{p_{\text{data}} + p_G} \right) \right) d\mathbf{x} \\ = & \min_G \int_{\mathbf{x}} \left(p_{\text{data}} \log \left(\frac{2p_{\text{data}}}{p_{\text{data}} + p_G} \right) + p_G \log \left(\frac{2p_G}{p_{\text{data}} + p_G} \right) \right) d\mathbf{x} - \log 4 \end{aligned}$$

Let us now find the optimal generator given the optimal discriminator.

$$\begin{aligned} & \min_G \int_{\mathbf{x}} \max_D (p_{\text{data}}(\mathbf{x}) \log D(\mathbf{x}) + p_G(\mathbf{x}) \log(1 - D(\mathbf{x}))) d\mathbf{x} \\ &= \min_G \int_{\mathbf{x}} \left(p_{\text{data}} \log \left(\frac{p_{\text{data}}}{\frac{p_{\text{data}}+p_G}{2}} \right) + p_G \log \left(\frac{p_G}{\frac{p_{\text{data}}+p_G}{2}} \right) \right) d\mathbf{x} \\ &= \min_G \left(D_{\text{KL}} \left(p_{\text{data}} \parallel \frac{p_{\text{data}} + p_G}{2} \right) + D_{\text{KL}} \left(p_G \parallel \frac{p_{\text{data}} + p_G}{2} \right) \right) \end{aligned}$$

Jensen-Shannon Divergence: Between two distributions f and g is symmetric and defined as

$$D_{JS}(f \| g) = \frac{1}{2} D_{\text{KL}} \left(f \parallel \frac{f+g}{2} \right) + \frac{1}{2} D_{\text{KL}} \left(g \parallel \frac{f+g}{2} \right)$$

The $D_{JS}(f \| g)$ is nonnegative and zero only if both the distributions are equal that is $f = g$. Therefore, the optimal generator $p_G = p_{\text{data}}$.

Generative Adversarial Networks

Jointly train generator G and discriminator D with a minimax game.

$$\min_G \max_D \left(\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(z)} [\log(1 - D(G(\mathbf{z})))] \right)$$

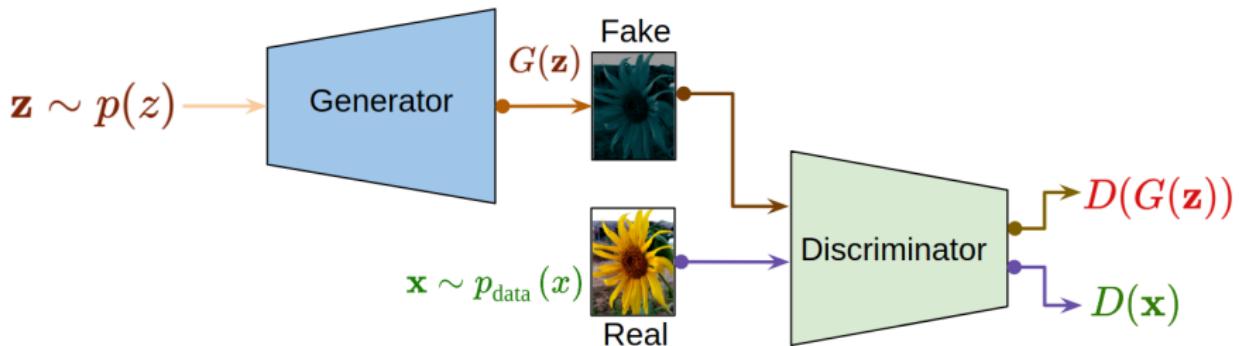
Claim: The global minimum of this game occurs at $p_G = p_{\text{data}}$.

Result: The global minimum of the minimax occurs, if

- $D^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$
- $p_G(\mathbf{x}) = p_{\text{data}}(\mathbf{x})$

Generative Adversarial Networks

Problem Statement: Given a dataset $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ of samples drawn from $p_{\text{data}}(\mathbf{x})$, draw new samples from p_{data} with explicitly modeling it.



$$\min_G \max_D \left(\underbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})]}_{D \text{ wants } D(\mathbf{x}) = 1 \text{ for real point}} + \underbrace{\mathbb{E}_{\mathbf{z} \sim p(z)} [\log(1 - D(G(\mathbf{z})))]}_{D \text{ wants } D(G(\mathbf{z})) = 0 \text{ for fake point}} \right)$$

The global minimum of the minimax occurs, if $D^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$ and $p_G = p_{\text{data}}$.

Algorithm 1 GAN Training

- 1: A dataset $\mathcal{S} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$.
- 2: Sample a batch of m training points $\{\mathbf{x}_{i_1}, \mathbf{x}_{i_2}, \dots, \mathbf{x}_{i_m}\}$ from \mathcal{S} .
- 3: Draw a batch of m latent vectors $\{\mathbf{z}_{i_1}, \mathbf{z}_{i_2}, \dots, \mathbf{z}_{i_m}\}$ from $p_{\mathbf{z}}$.
- 4: Update the parameters $\boldsymbol{\theta}_g$ of the generator $G_{\boldsymbol{\theta}_g}$ by SGD

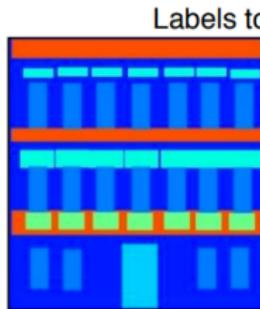
$$\begin{aligned}\nabla_{\boldsymbol{\theta}_g} V(\boldsymbol{\theta}_g, \boldsymbol{\theta}_d) &= \frac{1}{m} \nabla_{\boldsymbol{\theta}_g} \sum_{k=i_1}^{i_m} \log(1 - D_{\boldsymbol{\theta}_d}(G_{\boldsymbol{\theta}_g}(\mathbf{z}_k))) \\ \boldsymbol{\theta}_g &\leftarrow \boldsymbol{\theta}_g - \alpha_g \nabla_{\boldsymbol{\theta}_g}\end{aligned}$$

- 5: Update the parameters $\boldsymbol{\theta}_d$ of the discriminator $D_{\boldsymbol{\theta}_d}$ by SGD

$$\begin{aligned}\nabla_{\boldsymbol{\theta}_d} V(\boldsymbol{\theta}_g, \boldsymbol{\theta}_d) &= \frac{1}{m} \nabla_{\boldsymbol{\theta}_d} \sum_{k=i_1}^{i_m} \log(D_{\boldsymbol{\theta}_d}(\mathbf{x}_k)) + \log(1 - D_{\boldsymbol{\theta}_d}(G_{\boldsymbol{\theta}_g}(\mathbf{z}_k))) \\ \boldsymbol{\theta}_d &\leftarrow \boldsymbol{\theta}_d + \alpha_d \nabla_{\boldsymbol{\theta}_d}\end{aligned}$$

Conditional-GANs [Isola et al. 2017]

Image-to-image translation: We are given a set of paired images from two domains, X and Y . Can we translate from $X \rightarrow Y$.



input



output

BW to Color



input

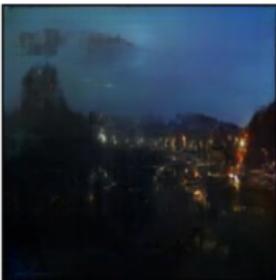


output

Day to Night



input



output

Edges to Photo



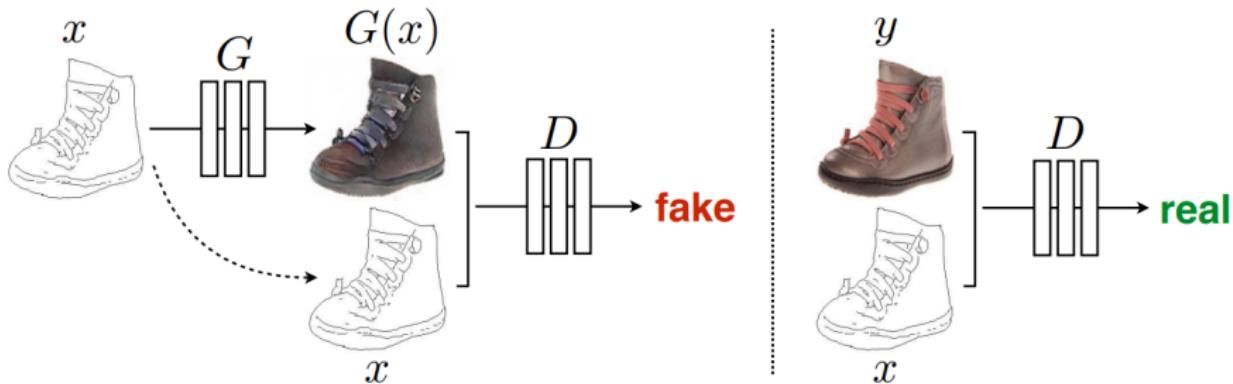
input



output

Conditional-GANs [Isola et al. 2017]

Image-to-image translation: We are given paired images from two domains, X and Y . Can we translate from $X \rightarrow Y$.



cGAN Losses

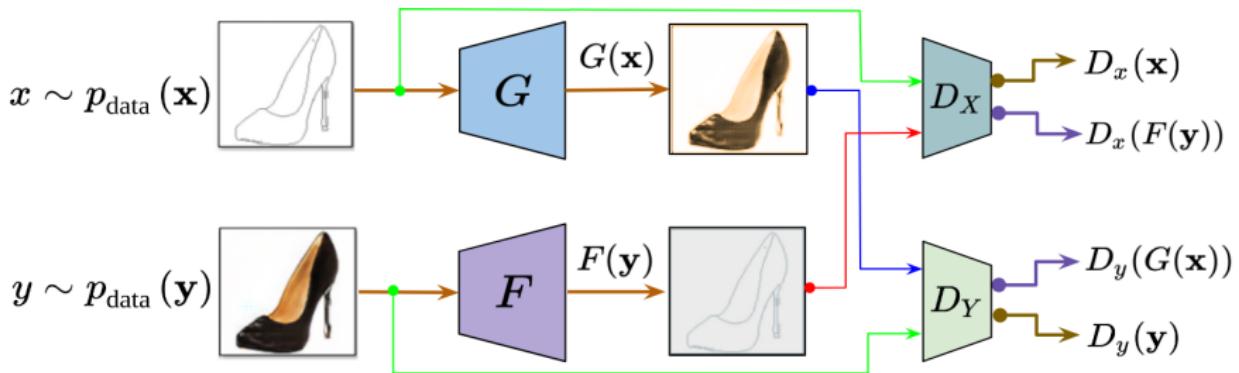
$$\ell_{\text{GAN}}(G, D) = \mathbb{E}_{\mathbf{x}, \mathbf{y}}[\log D(\mathbf{x}, \mathbf{y})] + \mathbb{E}_{\mathbf{x}}[\log(1 - D(\mathbf{x}, G(\mathbf{x})))]$$

$$\ell_c(G) = \mathbb{E}_{\mathbf{x}, \mathbf{y}}[\|\mathbf{y} - G(\mathbf{x}, \mathbf{z})\|_1]$$

$$\min_G \max_D \ell_{\text{GAN}}(G, D) + \lambda \ell_c(G)$$

Cycle-GANs [Zhu et al. 2020]

Image-to-image translation: We are given images from two domains, X and Y . Can we translate from $X \leftrightarrow Y$ in an unsupervised manner?



Adversarial Losses

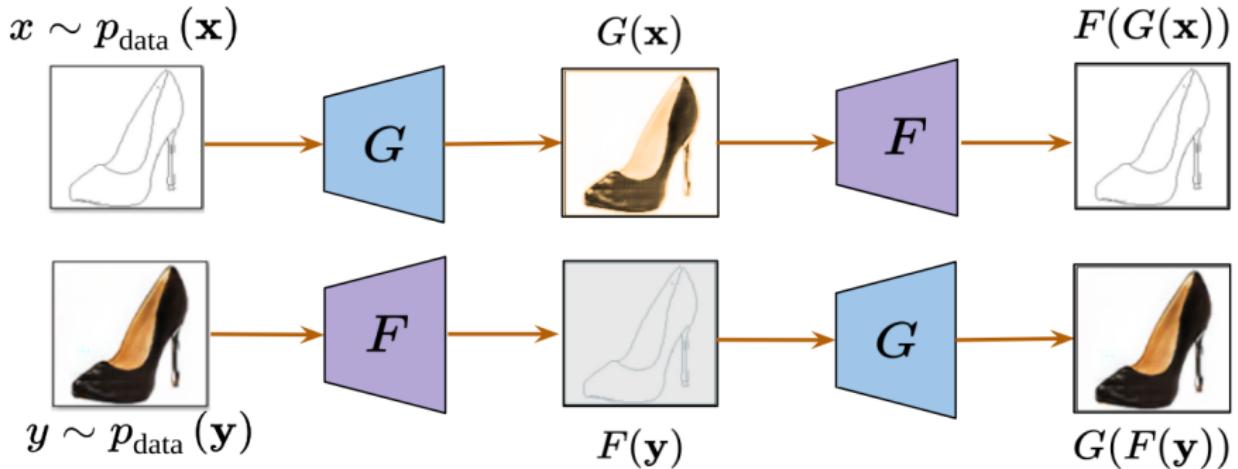
$$\min_F \max_{D_x} \ell_{\text{GAN}}(F, D_x)$$

$$\ell_{\text{GAN}}(F, D_x) = \mathbb{E}_{\mathbf{x}}[\log D_x(\mathbf{x})] + \mathbb{E}_{\mathbf{y}}[\log(1 - D_x(F(\mathbf{y})))]$$

$$\min_G \max_{D_y} \ell_{\text{GAN}}(G, D_y)$$

$$\ell_{\text{GAN}}(G, D_y) = \mathbb{E}_{\mathbf{y}}[\log D_y(\mathbf{y})] + \mathbb{E}_{\mathbf{x}}[\log(1 - D_y(G(\mathbf{x})))]$$

Cycle-GANs: Cycle Consistency



Cycle Consistency Losses

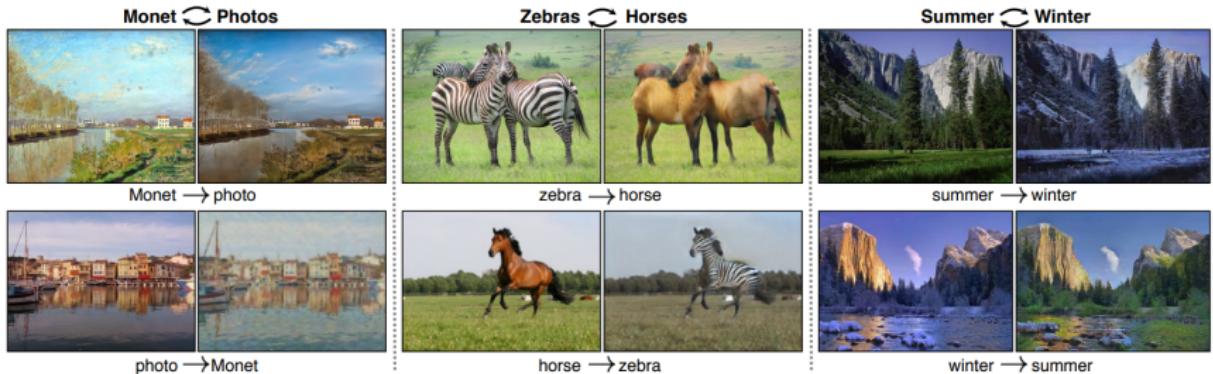
Learn the generators so that we can transform $X \leftrightarrow Y$ that is $\mathbf{x} = F(G(\mathbf{x}))$ and $\mathbf{y} = G(F(\mathbf{y}))$.

$$\ell_{\text{cyc}}(F, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\|\mathbf{x} - F(G(\mathbf{x}))\|_1] + \mathbb{E}_{\mathbf{y} \sim p_{\text{data}}(\mathbf{y})} [\|\mathbf{y} - G(F(\mathbf{y}))\|_1]$$

Cycle-GANs

Total Loss

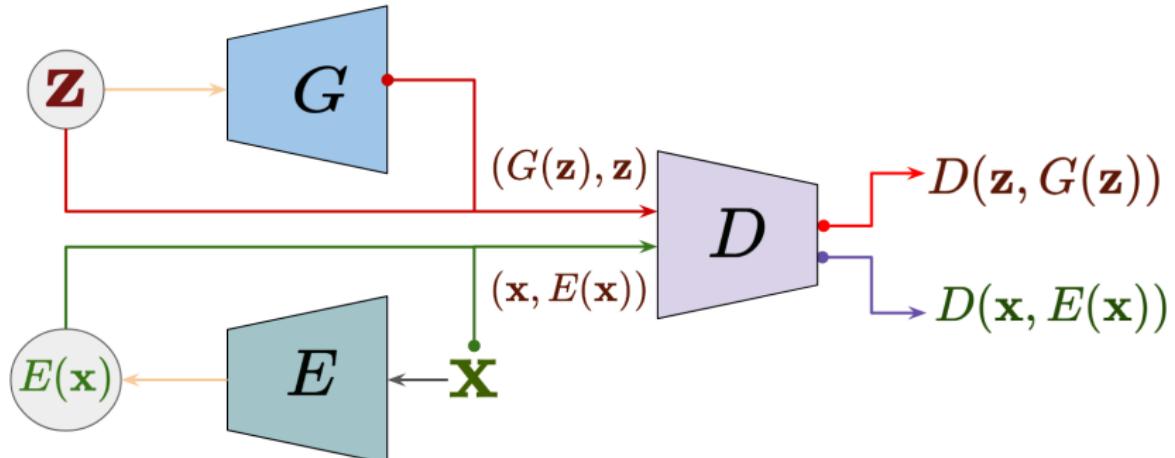
$$\ell(F, G, D_x, D_y) = \ell_{\text{GAN}}(F, D_x) + \ell_{\text{GAN}}(G, D_y) + \lambda \ell_{\text{cyc}}(F, G).$$



Understanding Latent Variables

- In variational auto-encoders, we can semantically control the latent vectors to generate new samples using the decoder.
- We can also use the encoder to find a semantic representation of an image from the learned distribution.
- These latent vectors can be used for feature representation of objects that can further be used for object classification.
- In GANs, given a latent vector \mathbf{z} , we can find the generated image using the generator as $G(\mathbf{z})$.
- However, given a new point from the learned distribution, can we map it to the corresponding latent representation \mathbf{z} ?
- In general, the generator is invertible.

Bidirectional GAN (BiGAN) [Donahue et al. 2017]



BiGAN Loss

$$\min_{G,E} \max_D (\mathbb{E}_{\mathbf{x}}[\log D(\mathbf{x}, E(\mathbf{x}))] + \mathbb{E}_{\mathbf{z}}[\log(1 - D(G(\mathbf{z}), \mathbf{z}))])$$

Claim: In order to “fool” a perfect discriminator D , BiGAN encoder E and generator G must invert each other. That is, $G(E(\mathbf{x})) = \mathbf{x}$ and $E(G(\mathbf{z})) = \mathbf{z}$.

Bidirectional GAN (BiGAN) [Donahue et al. 2017]

$G(\mathbf{z})$	
\mathbf{x}	
$G(E(\mathbf{x}))$	
$G(\mathbf{z})$	
\mathbf{x}	
$G(E(\mathbf{x}))$	
\mathbf{x}	
$G(E(\mathbf{x}))$	
\mathbf{x}	
$G(E(\mathbf{x}))$	

StyleGAN2 [Karras et al. 2020]



StyleGAN2 [Karras et al. 2020]

